

CONTENTS IN DETAIL

FOREWORD	xvii
-----------------	-------------

ACKNOWLEDGMENTS	xix
------------------------	------------

INTRODUCTION	xxi
---------------------	------------

What This Book Is About	xxii
Who This Book Is For	xxiii
AI Security vs. Cybersecurity vs. AI Safety	xxiv
This Book's Structure	xxv
The Python-Based Colab Notebooks	xxvii
Why Did I Write This Book?	xxviii
Let's Begin	xxix

PART I: AI AND SECURITY FUNDAMENTALS	1
---	----------

1 WHAT IS AI?	3
--------------------------	----------

Defining Machine Learning and the Origins of AI	4
The First Machine Learning Model: The Perceptron	5
An Example: Recommending a Movie	6
Neurons and Neural Networks	9
An Easy Introduction to the Math of Machine Learning	10
The Perceptron Formula Revisited	10
Data Structures	11
Matrices, Dataframes, and Tensors	12
Working with Models in Python	13
Processing Data	14
Creating Training and Test Sets	16
Building the Architecture	16
Compiling the Model	20
Training the Model	22
Evaluating the Model	23
Debunking Data Science Myths	25
Summary	26
Resources	26

2	WORKING WITH MODELS	29
Machine Learning Applications		30
Computer Vision		30
Signal Processing		32
Learning Methods		33
Supervised vs. Unsupervised Approaches		33
Reinforcement Learning		34
Performing Reinforcement Learning in Python		35
Generative AI		40
The Evolution of Large Language Models		40
The Embedding Space		42
Transformer Models		43
Reinforcement Learning with Human Feedback		46
Building a Language Model in Python		47
Creating Multimodal Models		52
Agents		58
Agent Communication Protocols		59
An Agent Communication Example		61
AI as a System		66
Retrieval-Augmented Generation		66
AI Infrastructure		69
Summary		72
Resources		73
 3	 AI THREATS	 75
The Attack Surface		76
Threat Modeling		77
System Diagrams		78
Example 1: MITRE ATLAS and OWASP		81
Example 2: MAESTRO		87
Threat Actors		90
Taxonomies		90
Criminal Enterprises		91
Advanced Persistent Threats		92
Threats to AI		92
AI Threat Intelligence		93
Traditional Software Vulnerabilities		94
AI-Enabled Cyber Crime		96
AI Weaknesses in Research		97
AI Weaknesses in the Real World		97
APT Interest in Intellectual Property		98
AI Security vs. Traditional Cybersecurity		99
Conceptual Models		100
Vulnerability Scoring		102
Uninterpretability		103
Unpatchability		105
Nondeterminism		105
Scale		105

Summary	106
Resources	106

PART II: ATTACKING AND DEFENDING AI

109

4

ATTACKS AND WEAKNESSES

111

The Machine Learning Life Cycle and Attack Surface	112
Disruption and Deception Methods.	114
Fast Gradient Sign Method and Projected Gradient Descent.	115
Adversarial Patch	122
Carlini and Wagner Attacks	126
Data Poisoning	129
Backdoors	130
Disclosure-Based Methods	133
Prompt Injection	133
Membership Inference	140
Model Extraction.	143
Side-Channel Attacks.	145
System-Level AI Threats	147
Retrieval Augmented Generation Attacks	148
Supply Chain Risks	148
Risks to Agents	148
Summary	153
Resources	154

5

DEFENSES, CONTROLS, AND MITIGATIONS

155

Machine Learning Defenses.	156
Adversarial Training	156
Gradient Obfuscation	159
Randomized Smoothing	163
Certified Defenses	166
Defensive Distillation	169
Model Ensembles	171
Statistical Detection Mechanisms.	173
System-Level Controls and Mitigations.	176
Input Validation.	176
Instruction Layer Filtering	178
Cybersecurity Controls and Mitigations	179
Agent-Specification Considerations	180
Security Across the Life Cycle	181
Mitigating Real-World Attacks	182
Facial Recognition Bypass (Without Liveness Detection)	183
Facial Recognition Bypass (With Liveness Detection)	183
Denial of Service in MathGPT.	184
Malicious Code in Google Colab Notebooks	185

Malware Detector Bypass.	185
Metamorphic Ransomware Variants.	186
Summary	187
Resources	187

PART III: THE AI SECURITY ECOSYSTEM 189

6 RED TEAMING AI 191

The Red Teaming Ecosystem	192
Cyber Red Teaming.	193
AI Red Teaming	193
Other Testing Methodologies	195
Planning an Exercise	196
Setting Objectives.	197
Defining the Threat Model	198
Selecting Adversarial Techniques	199
Choosing Tools and Frameworks.	201
Executing the Red Team Plan	203
Test Case 1: Prompt Injection	203
Test Case 2: Prompt Fuzzing.	206
Test Case 3: Multi-Turn Attacks	208
Responsible Disclosure and Reporting.	212
Coordinated Vulnerability Disclosure	212
Technical vs. Business Reports.	213
Report Examples.	213
Games and Challenges	218
Summary	219
Resources	220

7 ATTACKING AND DEFENDING WITH AI 221

AI for Cyber Defense	222
Anomaly and Intrusion Detection.	222
Malware Detection and Analysis.	225
Threat Intelligence.	227
Security Operations.	228
AI for Offensive Cybersecurity	230
Vulnerability Discovery and Exploitation.	231
Phishing and Social Engineering.	232
Autonomous Attack Agents.	235
Misinformation and Disinformation	237
Cyber Espionage and Surveillance	239
Summary	240
Resources	240

8

AI SAFETY

243

A History of AI Safety	244
Example: My AI Assistant's Possible Safety Issues	245
Short-Term Risks	247
Bias	247
Misuse	249
Interpretability	251
Job Displacement	254
Long-Term Risks	254
Alignment and the Control Problem	255
Failure Modes	256
Evaluations and Benchmarks	259
Testing for Harm	260
Classifying Toxicity with Detoxify	261
Evaluating Natural Language	262
Measuring Math Ability	264
Evaluating Reasoning	264
Rating Claude's Security Advice	267
Measuring Agent Safety	271
Mechanistic Interpretability	274
Activation Patching	274
Feature Attribution	275
Probing Classifiers	277
Circuit Analysis	278
Summary	279
Resources	279

9

AI GOVERNANCE

281

Two Approaches to AI Governance	282
Creating Dedicated AI Laws	282
Adapting Sector-Based Laws	284
Types of Guidelines	285
Policy	285
Standards	286
Industry Frameworks	287
Case Studies	289
The Evolution of Chatbot Regulations	290
The Complexities of Regulating Third Parties	291
Dead Rhinos and Responsibility Across Jurisdictions	291
Risk Analysis	292
Qualitative vs. Quantitative Approaches	293
Rapid Risk Audits	295
Bayesian Methods	299
Prioritization Matrixes	301

Implementing Governance in Practice	302
All Organizations	303
High-Risk Industries and AI Developers.	304
Summary	306
Resources	307

10

WHAT'S NEXT FOR AI SECURITY? 309

Technical Challenges	310
Agentic AI	310
Open vs. Closed Model Weights	312
Web3	314
Quantum Computing	315
Artificial General Intelligence	316
Social and Organizational Challenges	317
Model Development vs. Licensing	317
Geopolitics.	318
Communications and Design	319
Climate	321
The Workforce of the Future	322
Getting into AI Security.	323
Choosing Your Career Path	323
Honing Your Craft	324
Summary	326
Resources	327

CONCLUSION: A NEW KIND OF HACKER 329

FIGURE CREDITS 331

INDEX 333