

INDEX

A

- A2A (Agent-to-Agent) protocol, 59–60
- Aadhaar system, 239–240
- Access Now, 251
- accuracy, 51, 265
- ACSC (Australian Cyber Security Centre), 93, 324
- activation function, 17
- activation patching, 274–275
- activation strength, 139
- activation visualization, 252–253
- active cyber defense, 222
- Ada Lovelace Institute, 289
- Adaptive Moment Estimation (Adam) optimizer, 20, 22
- advanced persistent threats (APTs), 92, 98–99
- Advanced Threat Research team (McAfee), 186
- AdvBench, 260
- Adversarial AI Attacks, Mitigations, and Defense Strategies* (Sotiropoulos), 154
- adversarial machine learning attacks, 113
 - black box attacks, 144
 - targeted deception, 120–121
 - untargeted disruption, 118–120
 - disclosure-based methods
 - membership inference, 140–143
 - model extraction, 143–145
 - output leakage, 83
 - prompt injection, 133–140
 - side-channel attacks, 145–147
 - disruption and deception methods
 - adversarial patch, 122–126
 - backdoors, 130–133
 - Carlini and Wagner attacks, 126–129
 - data poisoning, 129–130
 - evasion methods, 114
 - fast gradient sign method and projected gradient descent, 115–121
 - targeted deception, 120–121
 - untargeted disruption, 118–120
 - examples, 115
 - model attacks, 83
 - model distillation, 143
 - model extraction, 143–145
 - techniques, 199–201
- adversarial machine learning defenses
 - certified defenses, 166–169
 - defensive distillation, 169–171
 - gradient obfuscation, 158–163
 - model ensembles, 171–173
 - randomized smoothing, 163–166
 - statistical detection mechanisms, 173–175
 - training, 156–158
- adversarial obfuscation, 114
- Adversarial Robustness Toolbox (ART), 154
- AEGIS, 260
- AgentBench, 266
- agentic AI, 231, 310–312
 - human interaction with, 312
- Agentic Red Teaming Guide* (Cloud Security Alliance), 220
- Agent Name Server (ANS), 310
- agents, 58–59
 - cards, 59
 - communication example, 61–66
 - communication protocols, 59–61
 - ecosystem, 88
 - frameworks, 87
 - measuring safety of, 271–274
 - risks to, 148–153
 - specification considerations, 180–181
- Agent-to-Agent (A2A) protocol, 59–60

- Agrawal, Ajay, 327
- AI (artificial intelligence), 5
 - bill of materials, 78
 - foundation models, 40, 87
 - guardrails, 47
 - origin of term, 4
 - as a system
 - AI infrastructure, 69–72
 - retrieval-augmented generation, 66–69
 - system-level threats
 - retrieval-augmented generation attacks, 148
 - risks to agents, 148–153
 - supply chain risks, 148
- AI Act, 249, 282–283
- AI Action Plan, 284
- AI and Technology (Our World in Data), 327
- AI compliance policy, 305
- AIDA (Artificial Intelligence and Data Act), 284
- AI Daily Brief, The* (Whittemore), 154
- AI-enabled cybercrime, 96–97
- AI for cyber defense
 - anomaly and intrusion detection, 222–224
 - malware detection and analysis, 225–226
 - security operations, 228–230
 - threat intelligence, 227–228
- AI for offensive cybersecurity
 - autonomous attack agents, 235–237
 - cyber espionage and surveillance, 239–240
 - misinformation and disinformation, 237–239
 - phishing and social engineering, 232–235
 - vulnerability discovery and exploitation, 231–232
- AI governance, 281–282
 - approaches to, 282–285
 - case studies, 289–292
 - implementation in practice, 302–306
 - of developers, 304–306
 - by organizations, 303–304
 - risk analysis
 - Bayesian methods, 299–301
 - prioritization matrixes, 301–302
 - qualitative vs. quantitative approaches, 293–294
 - rapid risk audits, 295–299
 - threats, 86–87
 - types of guidelines, 285–289
- AI incident response team, 306
- AI infrastructure, 69–72
- AI laws, 282–285
- AI Red Team (Microsoft), 97
- AI red teaming
 - disclosure and reporting
 - Coordinated Vulnerability Disclosure, 212–213
 - report examples, 213–218
 - technical vs. business reports, 213
- executing plan
 - multi-turn attacks, 208–212
 - prompt fuzzing, 206–208
 - prompt injection, 203–206
- games and challenges, 219
- part of governance, 305
- phases
 - analysis, 197
 - execution, 196
 - planning, 196
 - reconnaissance, 196
 - reporting, 197
- planning exercises, 196–197
 - choosing tools and frameworks, 201–203
 - defining threat model, 198–199
 - selecting adversarial techniques, 199–201
 - setting objectives, 197–198
 - red teaming ecosystem, 192–196
- AI risk assessments, 303. *See also* risk assessments
- AI Risk Management Framework (AI RMF), 287, 290, 292
- AI safety, xxiv, 243
 - AI security and cybersecurity vs., xxiv–xxv
 - evaluations and benchmarks, 259–260
 - agent safety, 271–274
 - Claude’s security advice, 267–271

- math ability, 264
 - natural language, 262–264
 - reasoning, 264–267
 - testing for harm, 260–261
 - toxicity, 261–262
- goal marginalization, 257–258
- history of, 244–245
- inner alignment, 255
- issues, 245–247
- mechanistic interpretability, 274–278
 - benchmarks, 273–274
- model drift, 181–182
- outer alignment, 255
- risks
 - alignment and control problem, 255–256
 - bias, 247–249
 - deceptive alignment, 258
 - failure modes, 256–259
 - interpretability, 251–253
 - job displacement, 254
 - misuse, 249–251
 - resource acquisition, 258
 - reward hacking, 35, 257
- x-risk, 245
- AI Safety Camp, 244–245
- AI Safety Summits, 251
- AI security, xxviii
 - AI safety and cybersecurity vs., xxiv–xxv
 - attacks and weaknesses, 82, 85, 97–98
 - getting into, 323–326
 - incident response, 306
 - mitigations, 99, 179–180
 - risk assessments, 303
 - threat modeling, 77
 - traditional cybersecurity vs., 99–100
 - conceptual models, 100–101
 - nondeterminism, 105
 - scale, 105–106
 - uninterpretability, 103–105
 - unpatchability, 105
 - vulnerability scoring, 102
- AI Security Bootcamp, 325
- AI Security Forum, 194
- AI Security Fundamentals Course (Farlow), 326, 327
- AI Security Initiative, 324
- AI Security Podcast, The*, 106
- AI threats
 - AI security vs. traditional
 - cybersecurity, 99–100
 - conceptual models, 100–101
 - nondeterminism, 105
 - scale, 105–106
 - uninterpretability, 103–105
 - unpatchability, 105
 - vulnerability scoring, 102
 - AI threat intelligence, 93–94
 - AI-enabled cybercrime, 96–97
 - AI weaknesses, 97–98
 - APT interest in intellectual property, 98–99
 - traditional software
 - vulnerabilities, 94–96
- attack surface, 76–77
- threat actors, 90–93
- threat modeling, 77
 - MAESTRO example, 87–90
 - MITRE ATLAS and OWASP example, 81–87
 - system diagrams, 78–81
- zero-click exploits, 76
- AI training, 304
- AI use policies, 303
- AI Village Red Team Challenge, 219
- AI Vulnerability Scoring System (AIVSS), 103, 107, 293
- AI weaknesses, 97–98
- AI web scrapers, 238
- AI worms, 84
- AIxCC (DARPA AI Cyber Challenge), 236–237
- AI X-risk Research Podcast, The (AXRP)* (Filan), 279
- Alan Turing Institute, 289
- algorithms, 8
- Algorithms to Live By* (Christian and Griffiths), 26
- AlgorithmWatch, 251
- alignment, 138, 246, 255–256
- Alignment Research Center (ARC), 324
- AlphaGo* (Moxie Pictures and Reel As Dirt), 26
- Alstott, Jeff, 107

- amateur attempts, 90
- Amazon Web Services (AWS), 71
- Amnesty International, 250
- Amodei, Dario, 245
- AMPS (Arithmetic with Multistep Problem Solving), 264
- analysis with AI, 225–226
- anchoring, 295
- Anderson, Ross, 154
- anomaly detection, 222–224
- Anonymous, 330
- ANS (Agent Name Server), 310
- Anthropic
 - career opportunities, 323
 - communication, 319
 - Frontier AI Red Teaming Framework, 202–203
 - MAESTRO and MCP, 88
 - Model Context Protocol, 59
 - model weights, 313
 - Responsible Scaling Policy and AI Safety Levels, 288
- anthropic package, 268
- Apache server, 195
- API keys, 61
- Apple, 320–321
- application layer risks, 86
- application programming interface (API), 47
- application-specific integrated circuits (ASICs), 71
- APT19 (Deep Panda), 92
- APT28 (Fancy Bear), 92
- APT38 (Lazarus Group), 233
- APTs (advanced persistent threats), 92, 98–99
- ARC (Alignment Research Center), 324
- architecture
 - building for model, 16–20
 - in infrastructure, 69
 - system diagrams of, 78–79
 - zero-trust, 179
- Arithmetic with Multistep Problem Solving (AMPS), 264
- Arnold, Tim, 188, 240
- ART (Adversarial Robustness Toolbox), 154
- artificial general intelligence, 316–317
- artificial intelligence. *See* AI
- Artificial Intelligence* (Mitchell), 27, 327
- Artificial Intelligence* (Russell), 245
- Artificial Intelligence and Data Act (AIDA), 284
- Artificial Intelligence as the New Hacker* (Valencia et al.), 241
- Artificial Intelligence Video Interview Act, 284
- artificial neural networks, 9
- arXiv, 97, 319
- ASD (Australian Signals Directorate), xxviii, 232, 324
- ASICs (application-specific integrated circuits), 71
- ASL (AI Safety Level) framework, 288
- ASR (attack success rate), 103
- Assange, Julian, 330
- Assembly Bills 602 and 730, 283
- assets, 77
- assumptions, 197
- Atari games, 34
- Atlas of AI* (Crawford), 327
- Attacker Knowledge Base, 95
- attacking and defending with AI
 - AI for cyber defense
 - anomaly and intrusion detection, 222–224
 - malware detection and analysis, 225–226
 - security operations, 228–230
 - threat intelligence, 227–228
 - AI for offensive cybersecurity
 - autonomous attack agents, 235–237
 - cyber espionage and surveillance, 239–240
 - misinformation and disinformation, 237–239
 - phishing and social engineering, 232–235
 - vulnerability discovery and exploitation, 231–232
- attack narrative, 213
- attacks and weaknesses
 - on data, 82
 - denial-of-service, 85

- machine learning life cycle and
 - attack surface, 112–114
- mitigating real-world attacks, 182–187
- on models, 83
- nondeterministic, 102
- system-level AI threats, 148–153
- weaknesses of AI, 97–98
- attack success rate (ASR), 103
- attack surface, 76–77, 112–114
- attention head, 43
- attention rollout, 276–277
- Australia, 284–285, 289
- Australian Cyber Security Centre (ACSC), 93, 324
- Australian Department of Defense, xxviii
- Australian Signals Directorate (ASD),
 - xxviii, 232, 324
- AutoModel, 53
- autonomous attack agents, 235–237
- AutoTokenizer, 53
- availability, 100
- AWS (Amazon Web Services), 71
- AWS S3, 195
- Azure AI, 292

B

- backdoors, 130–133
- backpropagation, 21
- backup, 180
- balanced data, 14
- Bar-on, Yogev, 107
- batch size, 23
- Bayes, Thomas, 300
- Bayesian methods, 299–301
- Bayes’ theorem, 300
- BBQ (Bias Benchmark for QA), 260
- BEC (business email compromise), 233
- benchmarks. *See* evaluations and benchmarks
 - benchmarks
- Bengio, Yoshua, 73, 245
- Benioff, Marc, 254
- Bernstein, Jake, 307
- BERT, 46, 53, 54
- Bertino, Elisa, 187
- Beyond the Imitation Game Benchmark (BIG-Bench Hard), 266
- bias, 246–249
- Bias Benchmark for QA (BBQ), 260

- bias factor, 5
- bias vector, 12
- Biden, Joe, 285
- Bilingual Evaluation Understudy (BLEU) score, 51, 263
- Bing Chat Data Pirate, 290
- black box attacks, 144
- Black Hat Python* (Seitz and Arnold),
 - 188, 240
- blockchain, 314
- BloombergGPT, 318
- BlueDot, 245
- blue teaming, 192
- Boileau, Kiwi A., 240
- Bolstering Online Transparency (BOT) Law, 283
- Bostrom, Nick, 245, 254, 256, 280
- botnets, 91
- bots, 283
- Boyle, Kip, 307
- BPE (Byte-Pair Encoding), 44
- Bradley, Henry A., 107
- Brazil, 284
- Breakout* game, 34
- brittle models, 156
- BSides events, 326
- Bug Bounty Bootcamp* (Li), 188
- bug bounty programs, 211
- Bugcrowd, 211
- bugs, 94–95
- Building an Enterprise Chatbot* (Singh, Ramasubramanian, and Shivam), 73
- Burp Suite scanner, 231
- business email compromise (BEC), 233
- business reports from red teams, 213
- Byte-Pair Encoding (BPE), 44

C

- C2 (command and control), 97
- calibration, 295
- California, 283, 291
- CalypsoAI, 323
- Canada, 284
- Capabilities for MachinE Learning (CaMeL) architecture, 181
- careers in AI security, 323–326
- Carlini and Wagner attacks, 126–129
- cascading goal drift, 194

- Center for AI and Digital Policy (CAIDP), 289
- Center for AI Safety (CAIS), 245, 324
- Center for Security and Emerging Technology (CSET), 324
- Central Intelligence Agency (CIA), 145–146
- central processing units (CPUs), 71
- Centre for the Governance of AI (GovAI), 245, 324
- certifiable robustness, 166
- certification from ISO, 287
- chains of techniques, 199–200
- Chatbot Arena, 259
- chatbot regulations, 290–291
- ChatGPT, xxviii, 40, 98
- checkpoints, 148
- Chen, Xiaofeng, 187
- Chevrolet, 98, 290
- Chierici, Alberto, 73
- China
 - AI laws, 284
 - AI-powered surveillance, 250
 - Deep Panda, 92
 - facial recognition, 239
 - geopolitics, 318
 - quantum computing, 315
 - responsible AI, 286
 - zero-day market, 93
- Christian, Brian, 26
- CIA (Central Intelligence Agency), 145–146
- CIA (confidentiality, integrity, and availability) triad, 100
- CI/CD (continuous integration and continuous delivery), 83
- CIFAR-10, 156, 159
- circuit analysis, 278
- CISA (Cybersecurity and Infrastructure Security Agency), 93, 324
- CivAI, 234
- Clark, Jack, 279
- Claude
 - context window, 46
 - model weights, 313
 - security advice from, 267–271
- clean loss, 167
- Clearview AI, 239, 250
- Cleverbot, 41
- CleverHans, 154
- climate, 321–322
- CLIP (Contrastive Language–Image Pretraining), 46
- closed source models, 137
- closed vs. open model weights, 312–313
- Cloudflare, 324
- Cloud Security Alliance (CSA), 87, 194, 288, 307
- Cluley, Graham, 106, 240
- clusters, 72
- code sandboxing, 184
- cognitive biases, 237
- color channels, 13
- command and control (C2), 97
- Common Crawl, 238
- Common Voice, 263
- Common Vulnerabilities and Exposures (CVE), 95
- Common Vulnerability Scoring System (CVSS), 102, 107, 293
- communication, 319–321
- COMP-128-1, 143
- COMPAS recidivism prediction tool, 247–248
- complexity (difficulty), 102
- compliance policy, AI, 305
- Computerphile (YouTube channel), 106
- computer vision, 30–32
- Compute Unified Device Architecture (CUDA), 123
- conceptual models, 100–101
- confidentiality, 100
- confidentiality, integrity, and availability (CIA) triad, 100
- constants, 11
- content moderation algorithms, 249
- contextual hijacking, 135
- context window, 46
- Conti, 91–92, 319
- continuous integration and continuous delivery (CI/CD), 83
- continuous verification, 179
- Contrastive Language–Image Pretraining (CLIP), 46
- control problem, 255–256

controls, 100. *See also* mitigations
least privilege access, 179
liveness detection, 183–184
system-level
agent-specification considerations, 180–181
cybersecurity controls and mitigations, 179–180
input validation, 176–178
instruction layer filtering, 178–179
security across life cycle, 181–182
convergence, 58
convolution, 31
convolutional neural network, 31
Coordinated Vulnerability Disclosure (CVD), 212–213
Copilot, 318
Cormen, Thomas H., 27
Cornell University, 97
corrigibility, 256
Counterfit tool, 145
Courville, Aaron, 73
CPUs (central processing units), 71
craft, honing, 324–326
Crawford, Kate, 248, 321, 327
credential stuffing, 231
criminal enterprises, 91–92
cross-attention, 45
CrowdStrike, 93, 226, 324
Crucible, 219
CSA (Cloud Security Alliance), 87, 194, 288, 307
CSA (Cyber Security Agency of Singapore), 93
CSET (Center for Security and Emerging Technology), 324
CTI (cyber threat intelligence), 93
CUDA (Compute Unified Device Architecture), 123
CVD (Coordinated Vulnerability Disclosure), 212–213
CVE (Common Vulnerabilities and Exposures), 95
CVSS (Common Vulnerability Scoring System), 102, 107, 293
CyBench, 273
cyber-capable institutions, operations by, 91
cybercrime, 75–76
AI-enabled, 96–97
syndicates, 90–91
cyber defense. *See* AI for cyber defense
cyber espionage, 239–240
Cyber Flag exercise, 193
Cyber Kill Chain, 81, 107
cyber red teaming, 192
Cyber Risk Management Podcast, The, 307
cybersecurity
AI safety vs., xxiv–xxv
AI security vs., xxiv–xxv, 99–100
conceptual models, 100–101
nondeterminism, 105
scale, 105–106
uninterpretability, 103–105
unpatchability, 105
vulnerability scoring, 102
anomaly detection, 222–224
backdoors, 130–133
bugs, 94–95
business email compromise (BEC), 233
controls and mitigations, 179–180
credential stuffing, 231
cyber threat intelligence, 93
data breaches, 294
fuzzing, 206
indicators of compromise, 99
intrusion detection systems, 223
least privilege access, 179
ransomware, 95, 186–187
social engineering, 232–235
spear phishing, 232
spyware, 76, 239
Cyber Security Agency of Singapore (CSA), 93
Cybersecurity and Infrastructure Security Agency (CISA), 93, 324
Cyber Security Centre, Australian (ACSC), 93, 324
Cyber Security Meets Machine Learning (ed. Chen, Susilo, and Bertino), 187
cyber threat intelligence (CTI), 93
Cylance, 185, 226

D

DAO hack, 314
dApps (decentralized applications), 314

- Darknet Diaries* podcast, 106
- Darktrace, 96
- dark web, 227
- DARPA AI Cyber Challenge (AIXCC), 236–237
- Dartmouth Conference Summer Research Project on Artificial Intelligence, 4
- data
 - and convergence, 58
 - processing, 14–16
- data attacks, 82
- data breaches, 294
- data cards, 15
- data centers, 71, 180, 321–322
- data engineering, 69
- dataframes, 12–13
- data generator, 79
- data operations in MAESTRO, 87
- data poisoning, 113, 129–130
- Data Protection Act, 284
- data science myths, 25–26
- datasets library, 48
- data structures, 11–12
- decentralized applications (dApps), 314
- deception attacks. *See* disruption and deception attacks
- deceptive alignment, 258
- decision boundary, 5
- decoders, 45
- deepfake-as-a-service, 233
- Deep Learning* (Goodfellow, Bengio, and Courville), 73
- deep learning models, 9
- DeepMind, Google, 323
- Deep Panda (APT19), 92
- DeepSeek, 47, 98, 318
- “Deep Synthesis” law, 286
- DEF CON, 60, 107, 192
- defending with AI. *See* attacking and defending with AI
- defense in depth, 176
- defenses, machine learning. *See* adversarial machine learning defenses
- Deisenroth, Marc Peter, 26
- denial-of-service attacks, 85, 184–185, 194
- dense vector representation, 44
- Department of Defense, 324
- Department of Industry, Science and Resources, 289
- “Dependency” comic (Munroe), 311
- dependency compromises, 83–84
- deployment infrastructure, 87
- deserialization, 96
- design, 319–321
- deterministic policy, 34
- Detoxify, 261–262
- Detoxify class, 261
- detoxify package, 261
- development and deployment pipeline, 80–81
- Dewey, Daniel, 245
- Diagon Alley, 227–228
- differential trust profiles, 201
- difficulty (complexity), 102
- diffusion models, 46
- disclosure, 101
- disinformation, 237–239
- distillation, 143, 169–171
- Domain Name Server (DNS), 310
- Donald Bren School of Information and Computer Sciences, 14
- downstreaming tasks, 40
- Dreadnode, 219, 232
- drones, 250
- dual-use technologies, 249
- dynamic analysis, 186

E

- echo chambers, 237
- economic denial-of-service attacks, 194
- edge-case testing, 199
- 80,000 Hours Podcast, The* (Wiblin), 280
- electricity, 321
- elements, 11
- EleutherAI, 259, 324
- ElevenLabs, 97
- ELIZA model, 41
- embedded watermarking, 313
- embedding generator, 79
- embedding space, 42–43
- Employment Equality Framework Directive, 284
- encoders, 45

- endpoint security, 180–181
- entanglement, 315
- episodes, 38
- epochs, 7
- epsilon value, 37
- EPSS (Exploit Prediction Scoring System), 231
- equations, 10
- Ethereum, 314
- Ethics Guidelines for Trustworthy AI, 285
- Ethics of AI, The* (Chierici), 73
- Euclidean distance, 128
- European AI Alliance, 286
- European Union, 282, 284–285, 315
- European Union AI Act, 249, 282–283
- Evals Security Guide (OpenAI), 268
- evaluations and benchmarks, 259–274
 - agent safety, 271–274
 - Claude’s security advice, 267–271
 - evaluating models, 23–25
 - in MAESTRO, 87
 - math ability, 264
 - natural language, 262–264
 - reasoning, 264–267
 - testing for harm, 260–261
 - toxicity, 261–262
- evasion methods, 114
- e-waste, 322
- examples
 - agent communication, 61–66
 - English-French translator, 47–52
 - fruitcake, 33–35
 - GridWorld, 35–40
 - movie recommendation, 6–9
 - predicting relevance of text to images, 52–58
 - retrieval-augmented generation, 66–69
 - wine classification model, 13
 - building architecture, 16–20
 - compiling, 20–22
 - creating training and test sets, 16
 - evaluating, 23–25
 - processing data, 14–16
 - training, 22–23
 - visualizing in Netron, 104

- executive voicepacks, 233
- Exodus Intelligence, 93
- explainability, 25, 251–253
- exploitation, 37, 231–232
- exploit maturity, 102
- Exploit Prediction Scoring System (EPSS), 231

F

- F1 score, 265
- Facebook, 319, 321
- facial recognition bypass, 183–184
- Fact Extraction and VERification (FEVER), 263
- Factor Analysis of Information Risk (FAIR), 294
- FactScore, 263
- failure modes, 256–259
- Faisal, A. Aldo, 26
- Fancy Bear (APT28), 92
- Farlow, Harriet, 106, 154, 327
- Fast.ai, 73, 327
- fast gradient sign method (FGSM), 115–118
 - targeted deception, 120–121
 - untargeted disruption, 118–120
- FBI (Federal Bureau of Investigation), 75
- feature attribution, 275–277
- Federal Trade Commission (FTC), 284
- feed-forward network (multilayer perceptron), 45
- FEVER (Fact Extraction and VERification), 263
- Filan, Daniel, 279
- filter matrix, 31
- fine-tuning, 52
- Flan-T5 model, 67–68
- Flask, 61
- floating-point operations per second (FLOPs), 283
- Forest Blizzard, 98
- format injection, 135
- formulas, 10
- Fortifying the Agentic Web* (Cloud Security Alliance), 307
- foundation models, 40, 87
- frameworks, 81, 201–203

FraudGPT, 233
 Freund, Jack, 307
 Friedman, Lex, 279
 Frontier AI Framework (Meta), 288
 Frontier AI Red Teaming Framework (Anthropic), 202–203
 frontier labs, 26
 Frontier Model Forum, 251
 frontier models, 26
 FTC (Federal Trade Commission), 284
 Fullpath, 98, 290
 function, 10
 future of hacking, 329–330
 Future of Life Institute, 279, 289
Future of Life Podcast, 279
 fuzzing, 206

G

G7, 251
 games for AI red teaming, 219
Gandalf game, 139–140, 219
 Gans, Joshua, 327
 Gaussian noise, 164
 Gemini, 46, 112
 General Data Protection Regulation (GDPR), 282
 generative AI, 40. *See also* AI
Generative AI Security (Huang et al.), 154
 Generative Pretrained Transformer. *See* GPT
 Generative Red Team challenge, 192
Genius Makers (Metz), 73
 geopolitics, 318–319
 Gibson, William, 107
 goal marginalization, 257–258
 Goldfarb, Avi, 327
 Good Ancestors, 289
 Goodfellow, Ian, 73, 115
 Google, 212, 251, 262, 288
 Google Cloud, 71
 Google Colab notebooks, xxvii, 185.
See also Python-based Colab notebooks
 Google DeepMind, 323
 Google Photos, 247
 Google-Proof Q&A (GPQA), 263
 GovAI (Centre for the Governance of AI), 245, 324

governance. *See* AI governance
 governance, risk, and compliance (GRC), 303
 GPT (Generative Pretrained Transformer), 40, 45
 GPT-2, 46, 136, 138–139
 GPT-3, 184
 GPT-4, 46, 318
 GPUs (graphics processing units), 71, 123
 Grade School Math 8K (GSM8K), 264
 gradient clipping, 161
 gradient descent, 21
 gradient obfuscation, 158–163
 gradients, 21
 gradient space diagram, 21
 Graduate-level benchmark, 263
 graph database, 71
 Gray, Patrick, 240
 gray box attacks, 144
 GRC (governance, risk, and compliance), 303
 GridWorld, 35–40
 GridWorld class, 37
 Griffiths, Tom, 26
 Grubb, Sam, 241
 GSM8K (Grade School Math 8K), 264
 guardrails, 47
 guidelines, 285–289
 Guidelines for Secure AI System Development, 287

H

Hacken, 232
 HackerOne, 211
 HackerPrompt, 140, 219
 HackTheBox, 219, 325
 hacktivism, 330
 HaluEval, 263
 hard fork, 314
 hard labels, 170
 harm, testing for, 260–261
 HarmfulQA, 260
 harmonic mean, 265
 HarrietHacks, 154, 326–327
Harry Potter and the Methods of Rationality (HPMOR; Yudkowsky), 244–245, 279
 HELM (Holistic Evaluation of Language Models), 259

- HiddenLayer, 94, 288, 323
- higher-dimensional structures, 12
- high-risk industries implementing governance, 304–306
- Hiregowdara, Siddharth, 234
- Hiroshima Process, 251
- Hluhluwe-iMfolozi Park, 97, 291
- Hoffman, Andrew, 187
- Hopper, Grace, 95
- Horizon Europe, 285
- How AI Works* (Kneusel), 27
- How Cybersecurity Really Works* (Grubb), 241
- How to Measure Anything in Cybersecurity Risk* (Hubbard and Seiersen), 295, 307
- Huang, Ken, 154
- Hubbard, Douglas W., 295, 307
- Hugging Face, 48, 212, 288
- Human Compatible* (Russell), 280
- human factor, 200
- human in the loop / human on the loop, 102–103
- Human Rights Act, 284
- Human Rights Watch, 250

I

- IBM 701, 4
- IBP (interval bound propagation), 166
- ICO (Information Commissioner’s Office), 289
- identity and access management (IAM), 179
- ID.me, 183, 291
- IDSs (intrusion detection systems), 223
- Illinois, 283–284
- ImageNet, 58, 117, 248
- image recognition, 30
- images, 13
- Imitation Game, The* (Black Bear Pictures and Bristol Automotive), 27
- Import AI* newsletter (Clark), 279
- INAI SI (International Network of AI Safety Institutes), 272
- incident response team, AI, 306
- India, 239–240
- indicators of compromise (IOCs), 99
- indirect attacks, 134
- induction heads, 278

- industry frameworks, 287–289
- Inference Impact metric, 293
- Information Commissioner’s Office (ICO), 289
- infrastructure, 69
 - threats to, 85–86
- inner alignment, 255
- input validation, 176–178
- insider threats, 90–91
- Inspect, 259, 267
- instruction layer filtering, 178–179
- instrumental convergence, 258–259
- integer overflows and underflows, 314
- integrated gradients, 276
- integrity, 100
- Intercode, 273, 279
- International Network of AI Safety Institutes (INAI SI), 272
- interpretability, 25, 251–253
- interval bound propagation (IBP), 166
- Introduction to Algorithms* (Cormen et al.), 27
- intrusion detection, 222–224
 - intrusion detection systems (IDSs), 223
- intrusions, 222
- IOCs (indicators of compromise), 99
- ISO (International Organization for Standardization), 286–287
 - ISO/IEC 27001, 286
 - ISO/IEC 42001, 286, 290, 292
 - ISO/IEC 42005, 286
- Israel, 250

J

- JailbreakBench, 260
- Jailbreak Chat, 140
- jailbreaking, 133
- job displacement by AI, 254
- Jobs, Steve, 320, 330
- Jones, Jack, 307
- jump host, 81
- jurisdictions, responsibility across, 291–292

K

- Kaggle dataset, 13
- Kargu-2 drone, 250

- Karpur, Ajay, 107
- Keras's Sequential model, 32
- kernel matrix, 31
- Khryseai
 - adversarial attack example, 111
 - adversarial training, 156
 - creating, 13
 - defenses example, 155
 - infrastructure of, 69–71
 - name origin, 13
 - safety issues, 245–247
 - using prioritization matrixes, 301–302
- Kneusel, Ronald T., 27
- knowledge distillation, 169
- Kohs, Greg, 26
- KPMG, 281, 307
- L**
 - L2 distance, 127–128
 - Lahav, Dan, 107
 - Lakera, 94, 139, 323
 - language models, 47–52
 - large language models (LLMs), 40–41, 193
 - Large Model Systems Organization (LMSYS), 259
 - LaTeX, 264
 - Lawfare Institute, 307
 - Lawfare Podcast, The*, 307
 - laws for AI, 282–285
 - Lazarus Group, 92, 233, 319
 - learning methods, 33–35
 - learning rate (step size), 7
 - Learn Prompting, 325
 - least privilege access, 179
 - LeftoverLocals, 148
 - legislation, 282
 - Lehe, Lewis, 73
 - LessWrong, 274
 - LessWrong* newsletter, 279
 - lethal autonomous weapons (LAWs), 250
 - Lewis, Patrick, 73
 - Lex Fridman Podcast*, 279
 - Liberator platform, 232
 - LibriSpeech, 263
 - Libya, 250
 - licensing vs. model development, 317–318
 - Life 3.0* (Tegmark), 280
 - life cycle of machine learning, 112–114
 - LIME (Local Interpretable Model-agnostic Explanations), 252
 - limitations, 197
 - linear function graph, 18
 - linear regression model, 18
 - Lipschitz regularization, 166
 - liveness detection, 183–184
 - LLM-as-a-judge, 171, 273
 - LLMs (large language models), 40–41, 193
 - Lloyds Bank, 97
 - LM Harness, 259
 - LMSYS (Large Model Systems Organization), 259
 - local minimum, 22
 - Lockheed Martin, 107
 - logging, 179–180, 222, 224
 - logic bomb, 226
 - loss event frequency, 294
 - loss function, 159
 - loss magnitude, 294
 - low-resource languages, 136
- M**
 - Machine Intelligence Research Institute (MIRI), 245
 - machine learning, 4–5. *See also* AI
 - applications, 30–33
 - BERT, 46, 53, 54
 - computer vision, 30–32
 - convolutional neural network, 31
 - cosine_similarity, 67
 - deep learning models, 9
 - embedding space, 42–43
 - epochs, 7
 - evolution of large language models, 40–41
 - explainability, 25, 251–253
 - F1 score, 265
 - feed-forward network (multilayer perceptron), 45
 - foundation models, 40, 87
 - GradientTape, 161
 - image recognition, 30
 - image_transform pipeline, 54–55

- language models, 47–52
 - LLMs, 40–41, 193
- learning rate, 7
- life cycle and attack surface of, 112–114
- math of, 10–13
 - L2 distance, 127–128
 - local minimum, 22
- `MultimodalModel` class, 53
- natural language processing
 - models, 41
- neurons and neural networks, 9, 17
- and origins of AI, 4–5
 - neurons and neural networks, 9
 - Perceptron, 5–6
 - recommending movies
 - example, 6–9
- reinforcement learning with human
 - feedback, 46–47
- `Seq2SeqTrainer` object, 50
- transformer models, 43–46
- unsupervised learning, 33–34
- working with models in Python, 13–25, 31, 161
- Machine Learning* (Ng), 27
- Machine Learning for Cybersecurity* (Omar), 240
- machine learning security operations (MLSecOps), 181
- MAESTRO, 87–90
- malware, 223, 225–226
- Malware Data Science* (Saxe and Sanders), 26
- malware detector bypass, 185–186
- Mandiant, 93
- Manhattan norm, 128
- MarianMT, 48
- Massive Multitask Language Understanding (MMLU), 266
- Massive Multitask Language Understanding – Pro (MMLU-Pro), 263
- massive open online courses (MOOCs), 326
- math ability, measuring, 264
- Mathematics Aptitude Test of Humans (MATH), 264
- Mathematics for Machine Learning* (Deisenroth, Faisal, and Ong), 26
- MathGPT, 184–185
- math of machine learning, 10–13
 - L2 distance, 127–128
 - local minimum, 22
- Matplotlib, 24, 296, 298
- matrices, 11–13
- maximum membership advantage, 265
- McAfee, 186
- McCarthy, John, 4
- McCulloch, Warren, 9
- MCP (Model Context Protocol), 59–60, 88–90
- Mean Opinion Score (MOS), 263
- Measuring and Managing Information Risk* (Freund and Jones), 307
- mechanistic interpretability, 105, 274
 - activation patching, 274–275
 - circuit analysis, 278
 - feature attribution, 275–277
 - probing classifiers, 277–278
- membership inference, 140–143
- membership-inference attacker, 265
- mesa objective, 255
- Meta, 212, 254, 288, 323
- Metame, 186
- metamorphic ransomware variants, 186–187
- Metasploit, 231
- metrics, 197
- Metz, Cade, 73
- Microsoft
 - AI-driven malware analysis, 226
 - AI Red Team, 97
 - AI security guidance, 287
 - AI threat intelligence, 94
 - career opportunities, 323
 - chatbots, 290
 - and Hluhlwe-iMfolozi Park, 292
- Microsoft Azure, 71
- Microsoft Guide for Securing the AI-Powered Enterprise*, 287
- Miles, Robert, 280
- Mileva Security Labs, xxix, 94
- Millennium Challenge, 192
- Ministry of Defense, 324
- Minsky, Marvin, 4

- MIRI (Machine Intelligence Research Institute), 245
- misinformation, 237–239
- misuse of AI, 249–251
- MIT AI Lab, 4
- Mitchell, Melanie, 27, 327
- mitigations, 99. *See also* controls
 - controls vs., 100
 - for real-world attacks, 182–187
 - system-level
 - agent-specification considerations, 180–181
 - cybersecurity controls and mitigations, 179–180
 - input validation, 176–178
 - instruction layer filtering, 178–179
 - security across life cycle, 181–182
- Mitnick, Kevin, 330
- MITRE ATLAS
 - for red team exercises, 203
 - threat modeling with, 81–82
 - data attacks, 82
 - infrastructure and runtime threats, 85–86
 - model attacks, 83
 - operator and governance threats, 86–87
 - output and application layer risks, 86
 - supply chain and dependency compromises, 83–84
- MITRE ATT&CK framework, 81–82, 107
- MITRE Corporation, 93–95, 324
- MLflow, 148
- MLSecOps (machine learning security operations), 181
- MMLU (Massive Multitask Language Understanding), 266
- MMLU-Pro (Massive Multitask Language Understanding – Pro), 263
- MNIST dataset, 169
- $m \times n$ matrix, 11
- MobileNetV2, 117–118, 123–124, 159
- Model Accessibility metric, 293
- model attacks, 83
- Model Context Protocol (MCP), 59–60, 88–90
- model distillation, 143
- model extraction, 143–145
- model poisoning attack, 113
- models, 8
 - agents, 58–59
 - communication example, 61–66
 - communication protocols, 59–61
- AI as system
 - infrastructure, 69–72
 - retrieval-augmented generation, 66–69
- compiling, 20–22
- development vs. licensing, 317–318
- ensembles, 171–173
- generative AI, 40
 - embedding space, 42–43
 - evolution of large language models, 40–41
 - reinforcement learning with human feedback, 46–47
 - transformer models, 43–46
- investigating internals of, 137–139
- language models in Python, 47–52
- learning methods, 33–35
- machine learning applications, 30–33
- multimodal, 52–58
- reinforcement learning in Python, 35–40
- weights, 312–313
- working in Python with, 13
 - building architecture, 16–20
 - compiling model, 20–22
 - creating training and test sets, 16
 - evaluating model, 23–25
 - gradient clipping, 161
 - gradient descent, 21
 - processing data, 14–16
 - training model, 22–23
- momentum technique, 22
- Monte Carlo modeling, 295–296
- MOOCs (massive open online courses), 326
- MOS (Mean Opinion Score), 263
- MOSNet, 263

- movie recommendation example, 6–9
- Moxie Pictures, 26
- MQ-9 Reaper, 250–251
- MT-Bench, 259
- Multi-Agent System Threat Modelling Guide (OWASP), 88–90, 107
- multi-agent systems, 87
- multi-class classification, 17
- multi-head self-attention, 45
- multilayer perceptron, 45
- multimodal models, 52–58
- multi-turn attacks, 208–212
- Munroe, Randall, 311
- My Health Records Act, 284

N

- Nanda, Neel, 280
- National Cyber Force, 232
- National Cyber Security Centre (NCSC), 93, 287, 324
- National Institute of Standards and Technology (NIST), 287, 289–290, 292, 324
- National Quantum Initiative, 315
- national security, 250
- National Security Agency (NSA), 239, 324
- National Vulnerability Database, 102
- natural language, 262–264
- natural language processing (NLP) models, 41
- Netron tool, 104
- neural fingerprinting, 174
- neural networks, 9, 17
- Neuromancer* (Gibson), 107
- neurons, 9, 17
- Nevo, Sella, 107
- New Generation Artificial Intelligence Governance Principles, 286
- New York City, 283
- Ng, Andrew, 27
- ngrok tool, 61–62
- Nmap tool, 231
- noise, 24
- noisiness, 103
- nondeterminism, 105
- nondeterministic attacks, 102
- non-differentiable layers, 159
- North Korea, 75, 92, 319
- Nuclei scanner, 231
- NumPy library, 12, 14, 296, 297
- NVIDIA, 260

O

- obfuscated code, 185
- obfuscation, 158–163, 225
- object detection, 30
- objective function, 258–259
- objectives, 197–198
- observability in MAESTRO, 87
- offensive cybersecurity. *See* AI for offensive cybersecurity
- Office for AI, 289
- Office of Qualifications and Examinations Regulation (Ofqual), 252
- Office of the Australian Information Commissioner (OAIC), 289
- off switch, 256
- Omar, Marwan, 240
- one-hot encoding, 16
- one-shot method, 116
- Ong, Cheng Soon, 26
- Onion Router, The (Tor), 227
- OpenAI
 - AI safety, 245
 - API key, 61
 - bug bounty program, 212
 - career opportunities, 323
 - communication, 319
 - Evals Security Guide, 268
 - model weights, 313
 - OpenAI Evals, 259
 - OpenAI Gym, 40
 - OpenAI Python SDK, 204
 - Preparedness Framework, 288
 - Red Teaming Framework, 202–203
 - reinforcement learning with human feedback, 47
- Open Philanthropy, 245, 325
- open source models, 138
- open vs. closed model weights, 312–313
- operational security, 200
- Operation Dream Job, 233
- operations, 10
 - by cyber-capable institutions, 91

- operator threats, 86–87
- optimization, 247
- optimizers, 20
- opus-books* dataset, 48
- orchestration, 59
 - framework, 71
- Ord, Toby, 280
- organizational challenges. *See* social and organizational challenges
- Our World in Data, 327
- outer alignment, 255
- out-of-scope items, 197
- output leakage, 83
- outputs, 197
- overfitting, 24
- OWASP
 - AI red teaming, 194
 - applying MAESTRO framework to MCP, 88–90
 - Artificial Intelligence Vulnerability Scoring System, 293
 - threat modeling with, 81–82
 - data attacks, 82
 - infrastructure and runtime threats, 85–86
 - model attacks, 83
 - operator and governance threats, 86–87
 - output and application layer risks, 86
 - supply chain and dependency compromises, 83–84

P

- packed malware, 225
- Paglen, Trevor, 248
- Palo Alto Networks, 97, 226, 324
- pandas library, 12, 14, 268
- Paperclip Maximizer, 246–247, 255
- parallelizable encoders, 46
- Partnership on AI, 251, 289
- Payment Card Industry Data Security Standard (PCI DSS), 195
- penetration testing, 195
- perception, 59
- Perceptron, 5–6, 10
- permission hijacks, 194
- persistence, 84
- Perspective API, 262
- PGD. *See* projected gradient descent
- phone phreakers, 330
- phishing, 232–235
- pickle module, 96
- PIL/Image, 53
- pipeline, 67
- Pitts, Walter, 9
- pixels, 13
- playbooks, 200
- poison rate, 130
- policy, 34, 282, 285–286
- Pong* game, 34
- pop-ups, 272
- Powell, Victor, 73
- Precipice, The* (Ord), 280
- Prediction Machines* (Agrawal, Gans, and Goldfarb), 327
- Preparedness Framework (OpenAI), 288
- prioritization matrixes, 301–302
- privileged sandboxes, 103
- privilege escalation, 85
- privileges, 102
- probing classifiers, 277–278
- process injection, 225, 226
- professional opportunistic efforts, 90
- projected gradient descent (PGD), 115–118
 - targeted deception, 120–121
 - untargeted disruption, 118–120
- Project Maven, 251
- prompt fuzzing, 206–208
- prompt injection, 133–136
 - contradictions in, 135
 - direct attacks, 134
 - executing red team plans, 203–206
 - format injection, 135
 - on *Gandalf*, 139–140
 - goals in, 134
 - indirect attacks, 134
 - investigating model internals, 137–139
 - multi-turn, 135
 - overriding instructions in, 134
 - performing, 136–137
 - persuasion in, 135
 - roleplaying in, 135

- self-referential prompts, 135
- token smuggling, 135
- ProPublica, 247–248
- Protect AI, 323
- protocol confusion attacks, 311
- provenance tracking, 302, 305
- purple teaming, 200
- PyPi (Python Package Index)
 - package, 148
- PyRIT (Python Risk Identification Toolkit), 201–202, 205–206, 208
- Python
 - adding convolutional layer to model, 32
 - adding recurrent neural network layer, 33
 - building language models in, 47–52
 - reinforcement learning in, 35–40
 - working with models in, 13–25, 31, 161
- Python-based Colab notebooks, xxvii–xxviii
 - adversarial-patch.ipynb*, 122–125
 - adversarial-training.ipynb*, 157–158
 - agentcomms.ipynb*, 61–65
 - agent-hack.ipynb*, 148–153
 - backdoor.ipynb*, 131–132
 - CarliniWagner.ipynb*, 127
 - certified-defenses.ipynb*, 166–168
 - datapointing.ipynb*, 130
 - defensive distillation.ipynb*, 169–170
 - Ensemble-Methods.ipynb*, 171–172
 - gradient-obfuscation.ipynb*, 159–162
 - GridWorld.ipynb*, 36–40
 - input-validation.ipynb*, 176–177
 - instruction-layer-filtering.ipynb*, 179
 - language-translation-model.ipynb*, 47–51
 - mechinterp.ipynb*, 275
 - membershipinference.ipynb*, 141
 - mini-RAG.ipynb*, 66
 - model-extraction.ipynb*, 144–145
 - montecarlo.ipynb*, 296–299
 - multi-modal-model.ipynb*, 53–57
 - PGD.ipynb*, 116–121
 - promptinjection.ipynb*, 136–137
 - randomized_smoothing.ipynb*, 163–165
 - red-team.ipynb*, 203

- security-eval.ipynb*, 268–271
- sidechannelattack.ipynb*, 146
- statistical-detection-methods.ipynb*, 173–175
- toxicity-eval.ipynb*, 260
- visualize-embedding-space.ipynb*, 42–43

PyTorch (torch) framework, 48, 53

Q

- Q-learning, 35
- Q-table, 37
- qualitative and quantitative risk analysis, 293–294
- quantum computing, 315–316
- Quantum Flagship program, 315
- Q-value, 35, 37–38

R

- RAG (retrieval-augmented generation), 66–69, 148
- RAGAS (Retrieval-Augmented Generation Assessment), 263
- Ramasubramanian, Karthik, 73
- RAND Corporation, 90, 251
- randomized smoothing, 163–166
- randomness, 159
- random numbers, 16
- random reset, 22
- RAND report, 318
- ransomware, 95, 186–187
- Rapid7, 95
- rapid risk audits, 295–299
- RealToxicityPrompts, 260
- ReaperAI, 236, 241
- reasoning, 59, 264–267
- Recall-Oriented Understudy for Gisting Evaluation (ROUGE), 263
- recovery, 180
- recurrent neural networks, 32
- Red Team (Zenko), 220
- red teaming ecosystem, 192–196.
 - See also* AI red teaming
- Red Teaming Framework (OpenAI), 202–203
- reentrancy attacks, 314
- reference fingerprint, 174
- regulation of third parties, 291

- reinforcement learning, 34–35
 - in Python, 35–40
 - Q-learning, 35
 - Q-table, 37
 - Q-value, 35, 37–38
- reinforcement learning with human
 - feedback (RLHF), 46–47
- ReLU (Rectified Linear Unit) function, 19–20
- Replit, 312
- reporting, 197, 213–218
- representative data, 14
- research, 97
- residual connections, 45
- resilience planning, 317
- ResNet-50, 53, 54
- resource acquisition, 258
- responsible AI, 285
- responsible parties, 304
- Responsible Scaling Policy
 - (Anthropic), 288
- REST API server, 61
- retrieval-augmented generation (RAG), 66–69, 148
- Retrieval-Augmented Generation
 - Assessment (RAGAS), 263
- retrieval poisoning, 148
- REvil, 91
- reward hacking, 35, 257
- rhinos example, 291–292
- Rhysider, Jack, 106
- risk analysis, 292–293
 - Bayesian methods, 299–301
 - prioritization matrixes, 301–302
 - qualitative vs. quantitative
 - approaches, 293–294
 - rapid risk audits, 295–299
- risk assessments, 305
 - AI Vulnerability Scoring System
 - (AIVSS), 103, 107, 293
 - CVSS (Common Vulnerability
 - Scoring System), 102, 107, 293
 - EPSS (Exploit Prediction Scoring
 - System), 231
 - Factor Analysis of Information Risk
 - (FAIR), 294
 - prioritization matrixes, 301–302
 - rapid risk audits, 295–299
- RISK Bench, 260
- risk frameworks, 282
 - AI Safety Level (ASL)
 - framework, 288
- risks
 - to agents, 148–153
 - AI safety long-term, 254–259
 - alignment and control problem, 255–256
 - failure modes, 256–259
 - AI safety short-term
 - bias, 247–249
 - interpretability, 251–253
 - job displacement, 254
 - misuse, 249–251
- Risky Business* podcast, 240
- RLHF (reinforcement learning with
 - human feedback), 46–47
- RNN layer, 33
- RobustBench, 188
- Robust Intelligence, 94, 288, 323
- robust loss, 167
- Rochester, Nathaniel, 4
- Rosenblat, Frank, 5
- ROUGE (Recall-Oriented Understudy
 - for Gisting Evaluation), 263
- RSA algorithm, 315
- rules of engagement, 197
- runtime, 85
 - threats, 85–86
- Russell, Stuart, 245, 256, 280
- Russia, 92–93, 250, 319

S

- safety. *See* AI safety
- Salesforce, 254
- saliency maps, 275–276
- SamSam ransomware, 95
- Sanders, Hillary, 26
- sanitized markdown, 180
- SANS, 219
- Saxe, Joshua, 26
- scaffolding, 180
- scalars, 11
- scale, 105–106
- scaling values, 15–16
- SDK (software development kit), 47
- search engine wars, 238–239

Secure AI Framework (SAIF), 288

Secure Shell (SSH), 81

security. *See also* AI security

- across life cycle, 181–182
- and compliance, 87
- operations, 228–230
- processes, 304
- securitized technologies, 313

security community, 305–306

Security Engineering (Anderson), 154

security information and event management (SIEM), 224

Security of Critical Infrastructure Act, 285

security orchestration, automation, and response (SOAR), 229

Seiersen, Richard, 295, 307

Seitz, Justin, 188, 240

self-preservation, 258

sequence-to-sequence model, 50

serialization, 96

session tokens, 76

Shannon, Claude, 5

shape of tensors, 13

SHapley Additive exPlanations (SHAP), 252

Shin Bet intelligence agency, 250

Shivam, Shrey, 73

Shodan scanner, 231

side-channel attacks, 145–147

SIEM (security information and event management), 224

sigma (Σ) symbol, 10

signal processing, 32–33

signatures, 223

Sims, The, 257–258

Singh, Abhishek, 73

single-blind testing, 201

SinglePromptAttack module, 202

single-step method, 116

Singularity Institute for Artificial Intelligence, 244

SIPRI (Stockholm International Peace Research Institute), 250

skeleton key attack, 130

sklearn, 14

Skylight, 185

Smashing Security podcast, 106, 240

smishing, 232

smoothing, 163–166

Snowden, Edward, 239

SOAR (security orchestration, automation, and response), 229

social and organizational challenges of AI security

- climate, 321–322
- communications and design, 319–321
- geopolitics, 318–319
- model development vs. licensing, 317–318
- workforce of future, 322–323

social engineering, 232–235

softened probabilities, 170

softmax activation function, 17

softmax layer, 45

software development kit (SDK), 47

Software Engineering Benchmark (SWE)-bench, 266

software patches, 94

software vulnerabilities, 94–96

SolarWinds, 319

Sophos, 226

Sotiropoulos, John, 154

South Africa, 97, 291

South Korea, 250

Soward, Emily, 60

Space Invaders game, 34

Spamouflage, 237

spear phishing, 232

specification gaming, 35

spyware, 76, 239

SQL injection, 86

SSH (Secure Shell), 81

standards, 282, 286–287

Stanford, 259

- HAI, 327

static analysis, 186

- code, 195

statistical detection mechanisms, 173–175

step function graph, 18

step size (learning rate), 7

StereoSet, 260

stochastic policy, 34

Stockholm International Peace Research Institute (SIPRI), 250

- STRIDE model, 81, 107
- Stuxnet, 92
- Superintelligence* (Bostrom), 245, 280
- superposition, 315
- supervised learning, 33–34
- supply chain, 83–84, 148
 - attacks, 199
- surveillance, 239–240, 250
- Susilo, Willy, 187
- SWE (Software Engineering Benchmark)-bench, 266
- symbolic execution, 195
- Synopsys, 148
- synthetic dataset, 167
- system diagrams, 78–81
- system documentation, 303–304
- system-level AI threats
 - retrieval augmented generation
 - attacks, 148
 - risks to agents, 148–153
 - supply chain risks, 148
- system-level controls and mitigations
 - agent-specification considerations, 180–181
 - cybersecurity controls and mitigations, 179–180
 - input validation, 176–178
 - instruction layer filtering, 178–179
 - security across life cycle, 181–182
- system prompts, 133
- Szegedy, Christian, 115, 154

T

- tactics, techniques, and procedures (TTPs), 77, 99
- targeted attacks, 120
- targeted deception, 120–121
- targets, 7
- Tay chatbot, 290
- technical challenges of AI security
 - agentic AI, 310–312
 - artificial general intelligence, 316–317
 - open vs. closed model weights, 312–313
 - quantum computing, 315–316
 - Web3, 314–315

- technical signals intelligence (TECHSIGINT), xxviii
- technical skills, 324–325
- Tegmark, Max, 245, 280
- temperature values, 170
- TEMPEST attack, 145
- temporal difference, 38
- TensorFlow, 12, 14
- tensor processing units (TPUs), 71
- tensors, 12–13
- TensorTrust, 140
- Term Frequency-Inverse Document Frequency (TF-IDF), 67
- test cases, 197
- test sets, 16
- Texas, 284
- Theriault, Carole, 106, 240
- third-party regulation, 291
- threat actors, 90–93
- threat intelligence, 227–228
 - motivations in, 99
- threat modeling, 77.
 - See also* AI threats
 - in AI red teaming, 197–199
 - MAESTRO example, 87–90
 - MITRE ATLAS and OWASP
 - example, 81–82
 - data attacks, 82
 - infrastructure and runtime
 - threats, 85–86
 - model attacks, 83
 - operator and governance threats, 86–87
 - output and application layer
 - risks, 86
 - supply chain and dependency
 - compromises, 83–84
 - system diagrams, 78–81
- threats to AI, 92–93. *See also* AI threats
- 3 D's model, 101, 114
- TikTok, 249
- timesteps, 33
- token smuggling, 135
- tools
 - AI risk assessments of, 303
 - COMPAS recidivism prediction, 247–248
 - Counterfit, 145

- Netron, 104
- ngrok, 61–62
- Nmap, 231
 - for red teaming, 197, 201–203
- TfidfVectorizer, 67
- Tor (The Onion Router), 227
- torch (PyTorch) framework, 48, 53
- torch.nn module, 53
- torchvision/models, 53
- torchvision/transforms, 53
- toxicity, 261–262
- toxic language, 260
- ToxiGen, 260
- TPM (trusted platform module), 181
- TPUs (tensor processing units), 71
- tqdm package, 268
- traditional cybersecurity vs. AI security, 99–100
 - conceptual models, 100–101
 - nondeterminism, 105
 - scale, 105–106
 - uninterpretability, 103–105
 - unpatchability, 105
 - vulnerability scoring, 102
- traditional machine learning
 - methods, 30
- traditional software vulnerabilities, 94–96
- training
 - adversarial, 156–158
 - models, 22–23
 - regarding AI, 304
 - sets and tests, 16
- Training Data Integrity metric, 293
- transferability, 102
- transformer models, 40, 43–46
- transformers library, 48
- Transparency in Frontier Artificial Intelligence Act, 283
- trust boundaries, 311–312
- trusted platform module (TPM), 181
- trusted regions, 181
- trust radius, 103
- Trust Statement (xAI), 288
- TruthfulQA, 262–263
- TryHackMe, 219, 326
- TTPs (tactics, techniques, and procedures), 77, 99

- Turing, Alan, 27
- Two Minute Papers, 154
- Tyldum, Morten, 27

U

- UBI (universal basic income), 254
- underfitting, 24–25
- uninterpretability, 103–105
- United Kingdom
 - AI-driven offensive cybersecurity, 232
 - AI laws, 284–285
 - AI-powered facial recognition, 239
 - career opportunities, 324
 - changing policy law in, 289
 - grading algorithm, 252
 - industry frameworks, 287
- United Nations, 250
- United States
 - AI-driven offensive cybersecurity, 232
 - AI laws, 283–285
 - career opportunities, 324
 - changing policy law in, 289
 - geopolitics, 318
 - industry frameworks, 287
 - misuse of AI, 250
 - quantum computing, 315
- unit vector, 12
- universal basic income (UBI), 254
- unpatchability, 105
- unsupervised learning, 33–34
- untargeted disruption, 118–120
- untrusted regions, 181
- US Cyber Command, 193
- use policies, AI, 303
- utility function, 258–259
- Uyghur population, 239

V

- Valencia, Leroy J., 241
- validation split, 23
- Vaswani, Ashish, 73
- vectors, 11
- vector store, 79
- Vice*, 97
- Virginia, 284
- VirusTotal, 223

- vishing, 232
- Volt Typhoon, 319
- vulnerabilities
 - discovery, 231–232
 - research, 195
 - scanning, 195
 - scoring, 102
- Vulnerabilities Equities Process (VEP), 93
- Vulnerability Reward Program (VRP), 212

W

- water and data centers, 321–322
- watermarking, embedded, 313
- weaknesses. *See* attacks and weaknesses
- Web Application Security* (Hoffman), 187
- WebArena, 266
- web iterations, 313–315
- web scraping, 238, 312
- weights, 5, 312–313
- wget command, 15
- white-box attacks, 144
- Whittemore, Nathaniel, 154
- Wiblin, Rob, 280
- wine classification model, 13
 - building architecture, 16–20
 - compiling, 20–22
 - creating training and test sets, 16

- evaluating, 23–25
- processing data, 14–16
- training, 22–23
- visualized in Netron, 104

- Winogender, 260
- workforce of the future, 322–323
- WormGPT, 233
- worms, AI, 84
- Wozniak, Steve, 330

X

- xAI, 288
- Xinjiang, 239
- XKCD, 311
- x-risk, 245

Y

- YouTube, 249
- Yudkowsky, Eliezer, 254, 279

Z

- Zelenskyy, Volodymyr, 237
- Zenko, Micah, 220
- ZeroBench, 266–267
- zero-click exploits, 76
- zero days, 93
- Zerodium, 93
- zero-trust architecture, 179
- Zeus malware, 225